

Week 5: Hyperparameter Optimization & Advanced Search Strategies

Heart Failure Survival Analysis

MDST Project

Winter 2026

Outline

- 1 Week 4 Recap
- 2 Hyperparameter Optimization Overview
- 3 Approach 1: GridSearchCV
- 4 Approach 2: Random Search with Optuna
- 5 Approach 3: Bayesian Optimization
- 6 Comparison & Strategy
- 7 Advanced: Pruning
- 8 Kaggle Competition
- 9 Summary

Quick Recap: Week 4 - Classification Models

Algorithms Trained:

- ① Logistic Regression
- ② Random Forest
- ③ Support Vector Machines
- ④ K-Nearest Neighbors

Key Concepts:

- Train/Test split with stratification
- Feature normalization
- Evaluation metrics (accuracy, precision, recall, F1, ROC-AUC)

This Week: Make these models even better by optimizing their hyperparameters!

From Basic Models to Optimized Models

Last Week: We trained models with default hyperparameters

Example - Random Forest:

- n_estimators = 100 (default)
- max_depth = None (default)
- Test accuracy = 73%

This Week: Find better settings!

- n_estimators = 50 (tuned)
- max_depth = 10 (tuned)
- Test accuracy = ??? (could be better!)

Question: How do we find these optimal values systematically?

What Are Hyperparameters?

Two Types of Parameters:

Model Parameters	Hyperparameters
Learned from data during training	Set before training
Weights in neural networks Coefficients in logistic regression	Learning rate, batch size Number of iterations

Random Forest Hyperparameters:

- **n_estimators:** Number of trees (more → better but slower)
- **max_depth:** Maximum tree depth (deeper → more complex)
- **min_samples_split:** Minimum samples to split a node
- **max_features:** Number of features to consider per split

The Problem: Too Many Combinations!

Scenario: Fine-tune Random Forest

Parameter Ranges:

- n_estimators: [20, 50, 100, 150, 200] (5 values)
- max_depth: [5, 10, 15, 20] (4 values)
- min_samples_split: [2, 5, 10] (3 values)
- max_features: ['sqrt', 'log2'] (2 values)

Total Combinations:

$$5 \times 4 \times 3 \times 2 = 120 \text{ models}$$

If each model takes 1 second to train:

$$120 \text{ seconds} = 2 \text{ minutes}$$

With more parameters?

$$10 \times 10 \times 10 \times 10 = 10,000 \text{ models!}$$

Solution: Use smart search strategies instead of trying everything!

GridSearchCV: Exhaustive Search

Strategy: Test all combinations in the parameter grid

Process:

- ① Define parameter grid (all combinations)
- ② For each combination:
 - Train model on training data
 - Evaluate with k-fold cross-validation
 - Record performance
- ③ Return best parameters

Code Pattern: `grid = GridSearchCV(model, param_grid, cv=5)`

```
grid.fit(X_train, y_train)
```

```
best_params = grid.best_params_
```

Pros: Exhaustive, guaranteed best in grid

Cons: Slow, exponential growth with parameters

Results from GridSearchCV

Parameter Grid:

- n_estimators: [20, 50, 100, 150, 200]
- max_depth: [None, 10, 20]
- min_samples_split: [2, 5, 7, 10, 13]
- max_features: ['sqrt', 'log2']

Total Models Evaluated: 150

Results:

- Best Validation Accuracy: **76.53%**
- Best Parameters: n_estimators=50, max_depth=10, min_samples_split=5, max_features='sqrt'
- Test Accuracy: **73.33%**

Observation: Works well, but 150 evaluations is expensive with larger models!

Random Search: Smarter Sampling

Strategy: Randomly sample from parameter space (not all combinations)

Key Idea:

- Instead of grid: define ranges (continuous or discrete)
- Sample randomly from these ranges
- With enough trials, likely to find good parameters
- Much faster than GridSearch

Why It Works:

- You don't need to evaluate every grid point
- Randomness helps explore the space
- Embarrassingly parallel (can run trials in parallel)

Optuna Advantage: Better API than GridSearch, cleaner code

What is Optuna?

- Lightweight hyperparameter optimization framework
- Supports multiple search strategies
- Easy to use, flexible parameter types

Basic Workflow:

1. Define objective function (returns metric to maximize)
2. Define parameter ranges within objective function
3. Create study with desired sampler
4. Optimize by running trials

Code:

```
study = optuna.create_study(sampler=RandomSampler())
study.optimize(objective_fn, n_trials=50)
best_params = study.best_params
```

Random Search Results (50 trials)

Setup:

- Sampler: RandomSampler (random exploration)
- Number of trials: 50 (4x less than GridSearch!)
- Same parameter ranges as GridSearch

Results:

- Best Validation Accuracy: ?
- Best Parameters: **varies each run**
- Test Accuracy: ?

Advantages:

- 3x faster than GridSearch (50 vs 150 models)
- Can handle continuous parameter ranges
- Good for exploratory tuning

Disadvantages:

- No learning from previous trials

Bayesian Optimization: Intelligent Search

Key Idea: Learn from previous trials to propose better parameters

How It Works:

- ① Start with random trials to explore
- ② Build probability model of objective function
- ③ Use model to propose promising parameters
- ④ Evaluate new parameters
- ⑤ Update model with new observation
- ⑥ Repeat steps 3-5

Advantage Over Random Search:

- **Exploitation:** Try parameters near previous good results
- **Exploration:** Still explore uncertain regions
- Uses all past information, not just random samples

Result: Better performance with fewer trials!

Tree-structured Parzen Estimator (TPE)

What is TPE?

- State-of-the-art Bayesian optimization algorithm
- Used by Optuna by default
- Very effective and fast

Key Property: “Exploitation vs Exploration”

Early Trials:

- Mostly explore randomly
- Build initial understanding

Later Trials:

- Focus on promising regions
- Still explore uncertain areas

Code: `study = optuna.create_study(sampler=TPESampler())
study.optimize(objective_fn, n_trials=50)`

Bayesian Optimization Results (50 trials)

Setup:

- Sampler: TPESampler (Bayesian + intelligent search)
- Number of trials: 50
- Same parameter ranges and objective function

Results:

- Best Validation Accuracy: ?
- Best Parameters: **likely better than random**
- Test Accuracy: ?

Key Observation:

- Same number of trials as Random Search (50)
- But Bayesian typically finds better parameters
- Because it learns from previous trials

Practical Impact: Get GridSearch quality in 1/3 the time!

Comparing All Three Approaches

Quick Comparison Table:

	GridSearch	Random	Bayesian
Trials	150	50	50
Speed	Slow	Medium	Fast
Learning	None	None	YES
Best For	Small grids	Exploration	Most cases

Efficiency: Bayesian > Random > GridSearch

Quality: Bayesian \geq GridSearch > Random

When to Use Each Approach

Use GridSearchCV When:

- Parameter space is small and discrete
- You have plenty of compute resources
- Parameters are well-understood

Use Random Search When:

- Quick exploration needed
- Easily parallelizable
- Simple implementation preferred

Use Bayesian Optimization (Optuna) When:

- **[RECOMMENDED FOR MOST PROJECTS]**
- Training models is expensive/slow
- Mix of continuous and discrete parameters
- Want best results with reasonable compute
- Working on real problems (not toy examples)

Pruning: Early Stopping

Idea: If a model looks bad halfway through, stop training it!

Example Scenario:

- Trial 1 (bad hyperparams): epoch 1 acc=45%, epoch 10 acc=46%
- **Prune this trial** instead of training all 100 epochs
- Save computation time!

Useful For:

- Neural networks (many epochs)
- Expensive models
- When training takes hours/days

Code Pattern: for epoch in range(100):

```
acc = train_epoch()  
trial.report(acc, epoch)  
if trial.should_prune():  
    raise optuna.TrialPruned()
```

Kaggle-Style Challenge

Your Task:

- ① Use GridSearch with different ML classifiers
- ② Find best hyperparameters for each
- ③ Compare test accuracy across models
- ④ Goal: Highest test accuracy wins!

Fixed Parameters:

- Train/Test split: 70/30, stratified, random_state=21
- Cross-validation: 5-fold
- Metric: accuracy

Models to Try:

- Logistic Regression
- Random Forest
- Support Vector Machines
- Gradient Boosting (new!)

Key Takeaways

Hyperparameters matter! Default parameters are rarely optimal.

Three Search Strategies:

- ① **GridSearchCV**: Exhaustive but slow
- ② **Random Search**: Faster, less learning
- ③ **Bayesian (Optuna)**: Smart and efficient [RECOMMENDED]

Practical Recommendations:

- Start with GridSearch for small spaces
- Use Optuna for anything larger/expensive
- Always validate on test set, never train set!
- Compare multiple models, not just one

Next Steps:

- Implement hyperparameter tuning
- Compare models in the competition
- Explore advanced techniques (ensembles, stacking)

Week 5: Hyperparameter Optimization

Key Resources:

- GridSearchCV docs: scikit-learn
- Optuna: <https://optuna.org>
- Week 5 notebook has working examples